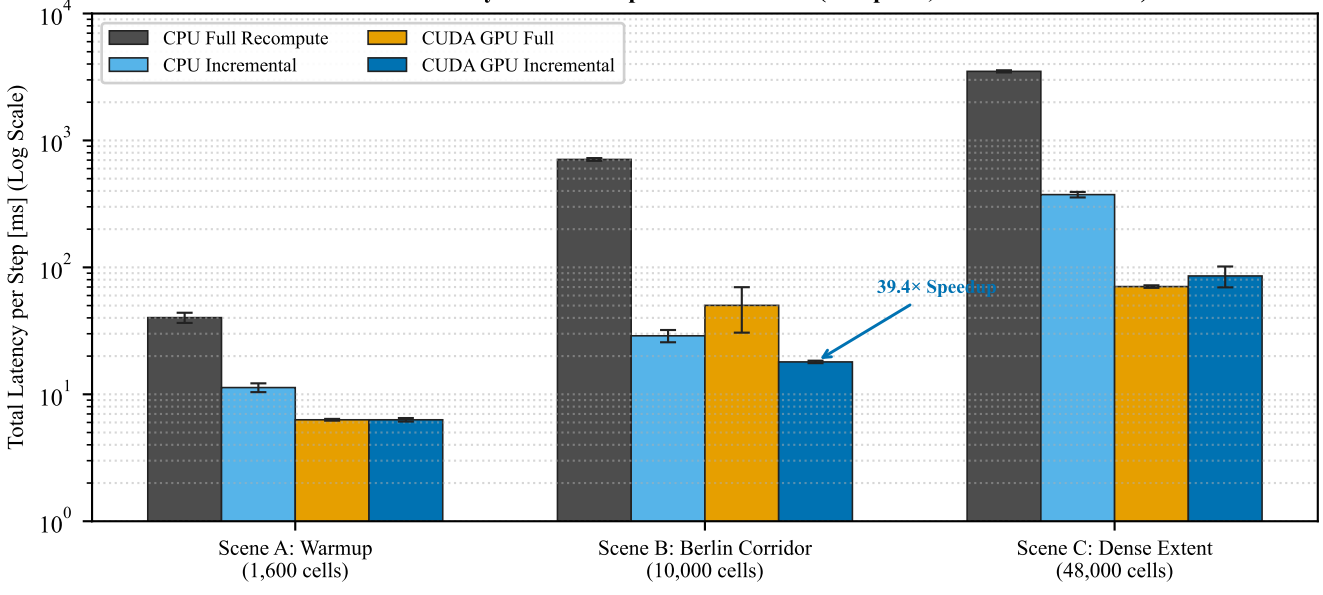
Dual-Backend Latency Across Computational Modes (5 Repeats, NVIDIA RTX 4050)
CPU Full Recompute
CPU Incremental
CUDA GPU Full
CUDA GPU Incremental
Total Latency per Step [ms] (Log Scale)
10^4
10^3
10^2
10^1
10^0
39.4× Speedup
Scene A: Warmup
(1,600 cells)
Scene B: Berlin Corridor
(10,000 cells)
Scene C: Dense Extent
(48,000 cells)